

# Summary of Reinforcement Learning: An Introduction

## Chapter 7: $n$ -step Bootstrapping

Justin Chen\*  
chenjus@bu.edu

Tyrone Hou\*  
tyroneh@bu.edu

January 22, 2018

### Introduction

- **Goals:**
  - Combine Monte Carlo methods and one-step temporal-difference (TD) methods
- **$n$ -step TD methods** - Generalization between both Monte Carlo methods and temporal-difference methods
- One-step TD methods - The same time step determines how often the action can be changed and the time interval over which bootstrapping is done
- In practice, we want to be able to quickly update the action value to account for changes, but bootstrapping performs over long timescales with significant changes
- $n$ -step methods enable bootstrapping over multiple steps

### 7.1 $n$ -step TD Prediction

- **What space of methods bridge the gap between Monte Carlo and TD methods?**
  - **Monte Carlo methods** estimate  $v_\pi$  for policy  $\pi$  using the *full return* from a given state.
  - **TD methods** bootstrap with *one-step* updates using current reward, but *proxy* the rest of the rewards using the current value estimate for the next state.
- **$n$ -step TD methods** do  *$n$ -step updates* with  $n$  actual rewards, proxy the rest of the return with  $v_\pi$  for the state  $n$  steps later. These generalize both Monte Carlo and one-step TD methods into one spectrum.
- **update target** - Quantity towards which a value estimate is moved during learning
 

Method	Target
Monte Carlo	$G_t \doteq R_{t+1} + \gamma R_{t+2} + \dots \gamma^{T-t-1} R_T$
one-step TD	$G_{t:t+1} \doteq R_{t+1} + \gamma V_t(S_{t+1})$
$n$ -step TD	$G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \dots \gamma^{n-1} R_{t+n} + \gamma^n V_{t+n-1}(S_{t+n})$

  - In any  $n$ -step methods, the target *approximates* the full return using  $n$  rewards. If  $t + n \geq T$  then  $G_{t:t+n} \doteq G_t$ .
- **$n$ -step return update rule**

$$V_{t+n}(S_t) \doteq V_{t+n-1}(S_t) + \alpha \left[ G_{t:t+n} - V_{t+n-1}(S_t) \right], \quad 0 \leq t < T \quad (1)$$

---

\*These authors contributed equally to this work

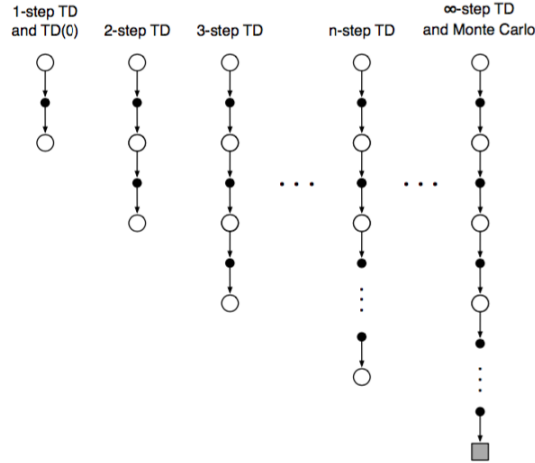


Figure 1: Backup diagrams for n-step methods

- Since this equation needs knowledge of rewards from  $n$  future steps, no updates can occur during the *first*  $n - 1$  steps of each episode.
- During the *last*  $n - 1$  steps the full return is always composed of less than  $n$  rewards, so we don't need to approximate using the value estimate. Instead we make updates using the full return.

#### n-step TD for estimating $V \approx v_\pi$

Initialize  $V(s)$  arbitrarily,  $\forall s \in \mathcal{S}$

Parameters: step size  $\alpha \in (0, 1]$ ,  $n \in \mathbb{Z}^+$

All store/access operations (for  $S_t$  and  $R_t$ ) can take their index mod  $n$

Repeat (for each episode):

Initialize and store  $S_0 \neq \text{terminal}$

$T \leftarrow \infty$

For  $t = 0, 1, 2, \dots$ :

If  $t < T$ , then:

Choose action using  $\pi(\cdot|S_t)$

Observe and store next state and reward pair  $S_{t+1}$  and  $R_{t+1}$

If  $S_{t+1}$  is terminal, then  $T \leftarrow t + 1$

$\tau \leftarrow t - n + 1$  (where an update occurs to the state estimate at time  $\tau$ )

If  $\tau \geq 0$ :

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

If  $\tau + n < T$ , then:  $G \leftarrow G + \gamma^n V(S_{\tau+n})$

$V(S_\tau) \leftarrow V(S_\tau) + \alpha[G - V(S_\tau)]$

Until  $\tau = T - 1$

## 7.2 n-step Sarsa

- How can  $n$ -step methods be used not just for prediction, but also for control?
- *n-step Sarsa* - Sarsa that utilizes bootstrapping over  $n$ -steps

$$Q_{t+n}(S_t, A_t) \doteq Q_{t+n-1}(S_t, A_t) + \alpha[G_{t:t+n} - Q_{t+n-1}(S_t, A_t)], 0 \leq t < T$$

- Values for all other states remain unchanged
- $Q_{t+n}(s, a) = Q_{t+n-1}(s, a)$  for all  $s \neq S_t$  or  $a \neq A_t$

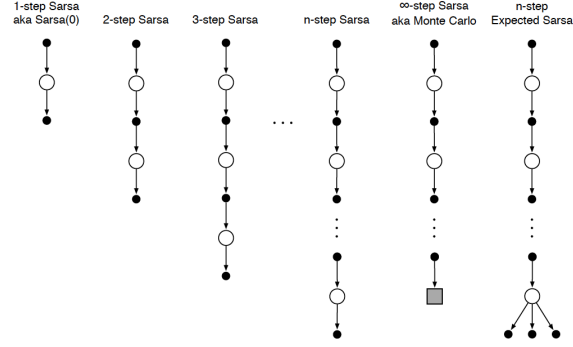


Figure 2: Backup diagrams for the spectrum of n-step methods for state-action values

### n-step Sarsa for estimating $Q \approx q_*$ , or $Q \approx q_\pi$ for a given $\pi$

Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}, a \in \mathcal{A}$

Initialize  $\pi$  to be  $\epsilon$ -greedy with respect to  $Q$ , or to a fixed given policy

Parameters: step size  $\alpha \in (0, 1]$ , small  $\epsilon > 0$ , a positive integer  $n$

All store and access operations (for  $S_t$ ,  $A_t$ , and  $R_t$ ) can take their index mod  $n$

Repeat (for each episode):

Initialize and store  $S_0 \neq \text{terminal}$

Select and store an action  $A_0 \tilde{\pi}(S_0)$

$T \leftarrow \infty$

For  $t = 0, 1, 2, \dots$ :

| If  $t < T$ , then:

| Take action  $A_t$

| Observe and store the next reward as  $R_{t+1}$  and the next state as  $S_{t+1}$

| If  $S_{t+1}$  is terminal, then:

|  $T \leftarrow t + 1$

| else:

| Select and store an action  $A_{t+1} \pi(S_{t+1})$

|  $\tau \leftarrow t - n + 1$  ( $\tau$  is the time whose estimate is being updated)

| If  $\tau \geq 0$ :

|  $G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

| If  $\tau + n < T$ , then  $G \leftarrow G + \gamma^n Q(S_{\tau+n})$

|  $Q(S_\tau, A_\tau) \leftarrow Q(A_\tau, S_\tau) + \alpha[G - Q(S_\tau, A_\tau)]$

| If  $\pi$  is being learned, then ensure that  $\pi(S_\tau)$  is  $\epsilon$ -greedy w.r.t.  $Q$

Until  $\tau = T - 1$

- Idea: switch states for actions or state-action pairs, and then proceed with  $\epsilon$ -greedy
- **n-step Returns** - Update targets that use  $n$  rewards over  $n$  time steps
- **n-step Expected Sarsa** is defined by the same equation as Sarsa, but the return is

$$G_{t:t+n} \doteq R_{t+1} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n \sum_a \pi(a|S_{t+n}) Q_{t+n-1}(S_{t+n}, a),$$

for all  $n$  and  $t$  such that  $n \geq 1$  and  $0 \leq t \leq T - n$

$$G_{t:t+n} \doteq R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q_{t+n-1}(S_{t+n}, A_{t+n}),$$

$n \geq 1, 0 \leq t < T - n$ , with  $G_{t:t+n} \doteq G_t$  if  $t + n \geq T$

## 7.3 $n$ -step Off-policy Learning by Importance Sampling

- Similar to importance sampling methods for Monte Carlo that compare target policy  $\pi$  and behavior policy  $b$ . Here we are only interested in the relative probabilities of the  $n$  actions taken in an  $n$ -step update.
- The product of the ratio of probabilities is called the **importance sampling ratio**  $\rho$ :

$$\rho_{t:t+h} \doteq \prod_{k=t}^{\min(h, T-1)} \frac{\pi(A_k|S_k)}{b(A_k|S_k)}$$

- If  $\pi$  never takes an action  $a$ , any  $n$  step return generated with action  $a$  is given zero weight.
- If  $\pi$  takes action  $a$  much more often than  $b$  does, the entire return is given a much larger weight.
- **Simple off-policy  $n$ -step TD update rule**

$$V_{t+n}(S_t) \doteq V_{t+n-1}(S_t) + \alpha \rho_{t:t+n-1} [G_{t:t+n} - V_{t+n-1}(S_t)], \quad 0 \leq t < T$$

- **off-policy  $n$ -step Sarsa update rule**

$$Q_{t+n}(S_t, A_t) \doteq Q_{t+n-1}(S_t, A_t) + \alpha \rho_{t+1:t+n-1} [G_{t:t+n} - Q_{t+n-1}(S_t, A_t)], \quad 0 \leq t < T$$

- $\rho$  starts at  $t + 1$  instead of  $t$ . We don't care about probability to take action  $A_t$ . We only want what happens to the return *given* we already selected the action.
- **off-policy  $n$ -step Expected Sarsa update rule**

$$Q_{t+n}(S_t, A_t) \doteq Q_{t+n-1}(S_t, A_t) + \alpha \rho_{t+1:t+n-2} [\mathbf{G}_{t:t+n} - Q_{t+n-1}(S_t, A_t)], \quad 0 \leq t < T$$

- Uses  $\rho_{t+1:t+n-2}$  instead of  $\rho_{t+1:t+n-1}$
- Expected Sarsa definition of expected return  $G_t$  (see 7.2)
- No need to importance sample *last time step* since we use the expected value and account for probability of taking each action under  $\pi$ .

**Off-policy  $n$ -step Sarsa for estimating  $Q \approx q_*$ , or  $Q \approx q_\pi$  for a given  $\pi$**

*Important changes in bold*

**Input: arbitrary  $\epsilon$ -soft behavior policy  $b$**

Initialize  $Q(s, a)$  arbitrarily,  $\forall s \in \mathcal{S}, a \in \mathcal{A}$

Initialize  $\pi$  as  $\epsilon$ -greedy w.r.t.  $Q$ , or as fixed policy

Parameters: step size  $\alpha \in (0, 1]$ , small  $\epsilon > 0$ ,  $n \in \mathbb{Z}^+$

All store/access operations (for  $S_t$ ,  $A_t$ , and  $R_t$ ) can take their index mod  $n$

Repeat (for each episode):

Initialize and store  $S_0 \neq$  terminal

Select and store  $A_0 \sim b(\cdot|S_0)$

$T \leftarrow \infty$

For  $t = 0, 1, 2, \dots$ :

If  $t < T$ , then:

Take action  $A_t$

Observe and store next state and reward pair  $S_{t+1}$  and  $R_{t+1}$

If  $S_{t+1}$  is terminal, then  $T \leftarrow t + 1$

else:

Select and store action  $A_{t+1} \sim b(\cdot|S_{t+1})$

$\tau \leftarrow t - n + 1$  (where an update occurs to the state estimate at time  $\tau$ )

If  $\tau \geq 0$ :

$$\rho \leftarrow \prod_{i=\tau+1}^{\min(\tau+n-1, T-1)} \frac{\pi(\mathbf{A}_i|\mathbf{S}_i)}{b(\mathbf{A}_i|\mathbf{S}_i)}$$

$$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$$

If  $\tau + n < T$ , then:  $G \leftarrow G + \gamma^n Q(S_{\tau+n}, A_{\tau+n})$

$$\mathbf{Q}(S_\tau, \mathbf{A}_\tau) \leftarrow \mathbf{Q}(S_\tau, \mathbf{A}_\tau) + \alpha \rho [\mathbf{G} - \mathbf{Q}(S_\tau, \mathbf{A}_\tau)]$$

If  $\pi$  is being learned, change  $\pi(\cdot|S_\tau)$  to be  $\epsilon$ -greedy w.r.t updated  $\mathbf{Q}$

Until  $\tau = T - 1$

## 7.4 \*Per-reward Off-policy Methods

- More efficient to use per-reward importance sampling
- Recursively defined n-step return:

$$G_{t:h} = R_{t+1} + \gamma G_{t+1:h}$$

- First reward and next state must be weight by importance sampling ratio at time  $t$  by  $\rho_t = \frac{\pi(A_t|S_t)}{b(A_t|S_t)}$
- Off-policy n-step return:

$$G_{t:h} \doteq \rho_t (R_{t+1} + \gamma G_{t+1:h}) + (1 - \rho_t) V_{h-1}(S_t), t < h \leq T, \text{ where } G_{t:t} \doteq V_{t-1}(S_t)$$

- Using policy  $b$  to try to model target  $\pi$
- If  $\pi$  would not select action at timestep  $t$ , then same as weighing  $b$  by  $\rho_t$  but this can cause high variance because of the resulting zeros
- Importance sampling ratio has expected value of one and is not correlated with the estimate

## 7.5 Off-policy Learning Without Importance Sampling: The $n$ -step Tree Backup Algorithm

- How to do off-policy learning without importance sampling? Extend the idea of n-step Expected Sarsa to create the **tree backup algorithm**
- The update equation backs up the current value estimate for state-action pair  $(S_t, A_t)$  using the action nodes of the tree below (the filled in circles).
- There are two 'half-steps' for the update, corresponding to two types of nodes:
  1. **Leaf nodes** - Accounts for estimated probability and Q-value of actions *not taken*. Each leaf node  $a$  contributes weight  $\pi(a|S_{t+j})$  to the target, the probability of action  $a$  occurring under target policy  $\pi$ .
  2. **Center nodes** - Uses reward to account for actions actually taken at each time step. Each center node  $A_{t+i}$  weights *all* the values from the next level with  $\pi(A_{t+i}|S_{t+i})$




Figure 3: 3-step tree backup update

- **One-step return target** is just Expected Sarsa:

$$G_{t:t+1} \doteq R_{t+1} + \gamma \sum_a \pi(a|S_{t+1})Q(S_{t+1}, A_{t+1}), \quad t < T - 1$$

$$G_{T-1:T} \doteq R_T$$

- **n-step return target** - Recursive definition with one-step return as base case.

$$G_{t:t+n} \doteq R_{t+1} + \gamma \sum_{a \neq A_{t+1}} \pi(a|S_{t+1})Q_{t+n-1}(S_{t+1}, a) + \gamma \pi(A_{t+1}|S_{t+1})G_{t+1:t+n}, \quad t+1 < T, n > 1$$

- **n-step tree backup update rule** - Same as n-step Sarsa

$$Q_{t+n}(S_t, A_t) = Q_{t+n-1}(S_t, A_t) + \alpha[G_{t:t+n} - Q_{t+n-1}(S_t, A_t)], \quad 0 \leq t < T$$

**n-step Tree Backup for estimating  $Q \approx q_\pi$ , or  $Q \approx q_\pi$  for a given  $\pi$**

Initialize  $Q(s, a)$  arbitrarily,  $\forall s \in \mathcal{S}, a \in \mathcal{A}$

Initialize  $\pi$  to be  $\epsilon$ -greedy w.r.t.  $Q$ , or as fixed policy

Parameters: step size  $\alpha \in (0, 1]$ , small  $\epsilon > 0$ ,  $n \in \mathbb{Z}^+$

All store/access operations can take their index mod  $n$

Repeat (for each episode):

Initialize and store  $S_0 \neq \text{terminal}$

Select and store  $A_0 \sim b(\cdot|S_0)$

Store  $Q(S_0, A_0)$  as  $Q_0$

$T \leftarrow \infty$

For  $t = 0, 1, 2, \dots$ :

If  $t < T$ , then:

Take action  $A_t$

Observe and store next state and reward pair  $S_{t+1}$  and  $R_{t+1}$

If  $S_{t+1}$  is terminal:

$T \leftarrow t + 1$

Store  $R - Q_t$  as  $\delta_t$

else:

Store  $R + \gamma \sum_a \pi(a|S_{t+1})Q(S_{t+1}, a) - Q_t$  as  $\delta_t$

Select arbitrarily and store action  $A_{t+1}$

Store  $Q(S_{t+1}, A_{t+1})$  as  $Q_{t+1}$

Store  $\pi(A_{t+1}|S_{t+1})$  as  $\pi_{t+1}$

$\tau \leftarrow t - n + 1$  (where an update occurs to the state estimate at time  $\tau$ )

If  $\tau \geq 0$ :

$Z \leftarrow 1$

$G \leftarrow Q_\tau$

For  $k = \tau, \dots, \min(\tau + n - 1, T - 1)$ :

$G \leftarrow G + Z\delta_k$

$Z \leftarrow \gamma Z\pi_{k+1}$

$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha[G - Q(S_\tau, A_\tau)]$

**If  $\pi$  is being learned, ensure  $\pi(\cdot|S_\tau)$  is  $\epsilon$ -greedy w.r.t updated  $Q(S_\tau, \cdot)$**

Until  $\tau = T - 1$

## 7.6 \*A Unifying Algorithm: $n$ -step $Q(\sigma)$

- Two ends of a spectrum: Sample action (Sarsa) or consider expectation over all actions (tree-backup algorithm)
- Expected Sarsa would sample all steps expect for the last step

$$\begin{array}{ccc} \text{Sarsa} & \longleftrightarrow & \text{Tree-backup} \\ \text{Always sample} & & \text{Never sample} \end{array}$$

- In between are methods that sample on some time steps and use expected transitions on others.
- Let each step have  $\sigma \in [0, 1]$ , a degree of sampling at time  $t$ .
- Sarsa has  $\sigma = 0$  for all time steps, Tree backup has  $\delta = 1$  for all steps.

$$\begin{array}{ccc} \sigma_t = 0 & \longleftrightarrow & \sigma_t = 1 \\ \text{Full sampling} & & \text{No sampling} \end{array}$$

- The  $n$ -step algorithm using  $\sigma$  is called  $\mathbf{Q}(\sigma)$




Figure 4: 4 types of  $n$ -step backup diagrams covered so far

- How do we define generalize  $n$ -step return using  $\delta$ ?
- Remember for  $n$ -step Sarsa

$$G_{t:t+n} = Q_{t-1}(S_t, A_t) + \sum_{k=t}^{\min(t+n-1, T-1)} \gamma^{k-t} [R_{k+1} + \gamma Q_k(S_{k+1}, A_{k+1}) - Q_{k-1}(S_k, A_k)]$$

- where we define  $\delta_k = R_{k+1} + \gamma Q_k(S_{k+1}, A_{k+1}) - Q_{k-1}(S_k, A_k)$  as the TD error.
- We can generalize TD error to 'slide' with  $\delta_k$  to take into account either more sample or expected values.

$$\begin{aligned} \delta_k &\doteq R_{k+1} + \gamma[\sigma_{t+1}Q_t(S_{t+1}, A_{t+1}) + (1 - \sigma_{t+1}\bar{Q}_{t+1})] - Q_{t-1}(S_t, A_t) \\ \bar{Q}_t &\doteq \sum_a \pi(a|S_t)Q_{t-1}(S_t, a) \end{aligned}$$

- Then we define  $n$ -step return for  $Q(\sigma)$  for the on-policy case as:

$$\begin{aligned} G_{t:t+1} &\doteq R_{t+1} + \gamma[\sigma_{t+1}Q_t(S_{t+1}, A_{t+1})(1 - \sigma_{t+1}\bar{Q}_{t+1})] &= \delta_t + Q_{t-1}(S_t, A_t) \\ G_{t:t+2} &\doteq R_{t+1} + \gamma[\sigma_{t+1}Q_t(S_{t+1}, A_{t+1}) + (1 - \sigma_{t+1}\bar{Q}_{t+1}) \\ &\quad - \gamma(1 - \sigma_{t+1}\pi(A_{t+1}|S_{t+1}))Q_t(S_{t+1}, A_{t+1}) \\ &\quad + \gamma(1 - \sigma_{t+1})\pi(A_{t+1}|S_{t+1})[R_{t+2} + \gamma[\sigma_{t+2}Q_t(S_{t+2}, A_{t+2}) + (1 - \sigma_{t+2}\bar{Q}_{t+2})]] \\ &\quad - \gamma\sigma_{t+1}Q_t(S_{t+1}, A_{t+1}) \\ &\quad + \gamma\sigma_{t+1}[R_{t+2} + \gamma[\sigma_{t+2}Q_t(S_{t+2}, A_{t+2}) + (1 - \sigma_{t+2}\bar{Q}_{t+2})]] \\ &= Q_{t-1}(S_t, A_t) + \delta_t \\ &\quad + \gamma(1 - \sigma_{t+1})\pi(A_{t+1}|S_{t+1})\delta_{t+1} \\ &\quad + \gamma\sigma_{t+1}\delta_{t+1} \\ &= Q_{t-1}(S_t, A_t) + \delta_t + \gamma[(1 - \sigma_{t+1})\pi(A_{t+1}|S_{t+1}) + \sigma_{t+1}]\delta_{t+1} \\ G_{t:t+n} &\doteq Q_{t-1}(S_t, A_t) + \sum_{k=t}^{\min(t+n-1, T-1)} \delta_k \prod_{i=t+1}^k \gamma[(1 - \sigma_i)\pi(A_i|S_i) + \sigma_i] \end{aligned}$$

- For the off-policy case, we use importance sampling and redefine  $\rho$  in terms of  $\sigma$ :

$$\rho_{t:h} \doteq \prod_{k=t}^{\min(h, T-1)} \left( \sigma_k \frac{\pi(A_k|S_k)}{b(A_k|S_k)} + 1 - \sigma_k \right)$$

**Off-policy n-step  $Q(\sigma)$  for estimating  $Q \approx q_\pi$ , or  $Q \approx q_\pi$  for a given  $\pi$**

Input: an arbitrary  $\epsilon$ -soft behavior policy  $b$

Initialize  $Q(s, a)$  arbitrarily,  $\forall s \in \mathcal{S}, a \in \mathcal{A}$

Initialize  $\pi$  to be  $\epsilon$ -greedy w.r.t.  $Q$ , or as fixed policy

Parameters: step size  $\alpha \in (0, 1]$ , small  $\epsilon > 0$ ,  $n \in \mathbb{Z}^+$

All store/access operations can take their index mod  $n$

Repeat (for each episode):

Initialize and store  $S_0 \neq \text{terminal}$

Select and store  $A_0 \sim b(\cdot|S_0)$

Store  $Q(S_0, A_0)$  as  $Q_0$

$T \leftarrow \infty$

For  $t = 0, 1, 2, \dots$ :

If  $t < T$ :

Take action  $A_t$

Observe and store next state and reward pair  $S_{t+1}$  and  $R_{t+1}$

If  $S_{t+1}$  is terminal:

$T \leftarrow t + 1$

Store  $R - Q_t$  as  $\delta_t$

else:

Select and store action  $A_{t+1} \sim b(\cdot|S_{t+1})$

Select and store  $\sigma_{t+1}$

Store  $Q(S_{t+1}, A_{t+1})$  as  $Q_{t+1}$

Store  $R + \gamma\sigma_{t+1}Q_{t+1} + \gamma(1 - \sigma_{t+1}) \sum_a \pi(a|S_{t+1})Q(S_{t+1}, a) - Q_t$  as  $\delta_t$

Store  $\pi(A_{t+1}|S_{t+1})$  as  $\pi_{t+1}$

Store  $\frac{\pi(A_t|S_t)}{b(A_t|S_t)}$  as  $\rho_{t+1}$

$\tau \leftarrow t - n + 1$  (where an update occurs to the state estimate at time  $\tau$ )

If  $\tau \geq 0$ :

$\rho \leftarrow 1$

$Z \leftarrow 1$

$G \leftarrow Q_\tau$

For  $k = \tau, \dots, \min(\tau + n - 1, T - 1)$ :

$G \leftarrow G + Z\delta_k$

$Z \leftarrow \gamma Z[(1 - \sigma_{k+1})\pi_{k+1} + \sigma_{k+1}]$

$\rho \leftarrow \rho(1 - \sigma_k + \sigma_k\rho_k)$

$Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha\rho[G - Q(S_\tau, A_\tau)]$

**If  $\pi$  is being learned, ensure  $\pi(\cdot|S_\tau)$  is  $\epsilon$ -greedy w.r.t updated  $Q(S_\tau, \cdot)$**

Until  $\tau = T - 1$

## 7.7 Summary

- This chapter focused a spectrum of temporal-difference methods with one-step TD at one end and Monte Carlo methods on the other
- Bootstrapping performs better than either extreme
- n-step methods look ahead at the next  $n$  rewards, states, and actions before updating which requires *more time and space per timestep* than methods that do not look ahead
- Increased cost of n-steps is an inherent limitation
- Advantage of n-step methods is that they're conceptually clear



## References

Sutton, Richard S., and Andrew G. Barto. “Temporal-Difference Learning.” Reinforcement Learning: An Introduction, The MIT Press, 2018, pp. 97-114.